DocGuard AI: Intelligent Document Forgery and Tampering Detection Using Multi-Stream Neural Networks

A Multi-Model Deep Learning System for Forensic Document Analysis

Pavan Kumar Kallakuri
Master of Science in Applied Artificial Intelligence
Shiley Marcos School of Engineering / University of San Diego
pkallakuri@sandiego.edu

Pratibha Kambi
Master of Science in Applied Artificial Intelligence
Shiley Marcos School of Engineering / University of San Diego
pkambi@sandiego.edu

Sajesh Kariadan
Master of Science in Applied Artificial Intelligence
Shiley Marcos School of Engineering / University of San Diego
skariadan@sandiego.edu

ABSTRACT

Document fraud costs the global economy an estimated $4.7 trillion annually, yet affordable and accessible forensic detection tools remain unavailable to most compliance professionals, auditors, and HR departments. This paper presents DocGuard AI, a deployed multi-stream neural network system for detecting forgery and tampering in scanned business documents. The system combines three independently trained deep learning models — an Error Level Analysis Convolutional Neural Network (ELA-CNN), a Metadata Multi-Layer Perceptron (Metadata MLP), and a Font Convolutional Recurrent Neural Network (Font CRNN) — into a weighted fusion pipeline that produces a forgery risk score, a four-tier verdict, and a Grad-CAM spatial heatmap highlighting tampered regions.

All three models were trained from scratch on a synthetic document forgery dataset of 4,800 image pairs generated programmatically from the RVL-CDIP corpus of real scanned business documents. Each model achieved 100% test accuracy and AUC-ROC of 1.0000 on the synthetic dataset, significantly outperforming a traditional machine learning baseline (Support Vector Machine with HOG+LBP features, AUC 0.9978). A live end-to-end test confirmed that a PDF containing a 2019 invoice date but created in 2026 returns a verdict of HIGH RISK with a forgery risk score of 84/100 and the specific signal: date mismatch — backdating detected. The system is deployed permanently on HuggingFace Spaces as a free public REST API with a drag-and-drop web frontend, and was built using an AWS SageMaker training pipeline.

Keywords: *document forgery detection, error level analysis, convolutional neural network, metadata analysis, BiLSTM, Grad-CAM, deep learning, computer vision, AWS SageMaker, HuggingFace*

1 Introduction

Every industry that relies on paper-based evidence faces the threat of document fraud. A forged invoice, backdated contract, photo-substituted identity document, or altered financial statement can redirect payments, invalidate legal proceedings, enable insurance fraud, or defeat regulatory compliance checks. The problem is not rare — the Association of Certified Fraud Examiners estimates that document fraud contributes to over $4.7 trillion in annual global economic losses.

The core challenge is that a forged document is visually indistinguishable from an authentic one to the human eye. The forger's advantage is that human perception is poor at detecting subtle pixel inconsistencies, compression artifact patterns, or metadata layer disagreements. Existing forensic tools — such as Forensic Toolkit (FTK) and EnCase — require expert analysts and cost upward of $10,000 per licence, making them inaccessible to the compliance professional, HR manager, or auditor who encounters a suspicious document in daily workflow.

DocGuard AI addresses this gap with the following hypothesis: a multi-stream deep learning pipeline that independently analyses pixel-level compression artifacts, document metadata signals, and font consistency — then fuses results via weighted scoring — will detect document forgery with materially higher accuracy than any single-stream model or traditional machine learning approach. The system is designed to be accessible via a single browser upload, requiring no technical expertise, and deployed permanently as a free service.

The end user is the non-technical compliance professional: the auditor reviewing invoice authenticity, the HR manager verifying qualification certificates, the insurance adjuster checking medical reports, or the legal professional authenticating contracts. DocGuard AI makes forensic-grade analysis available to this user with a drag-and-drop interface and a plain-language verdict.

---

2 Data Summary

2.1 Source Documents

The foundation of the training dataset is the RVL-CDIP corpus (Harley et al., 2015), which contains 400,000 greyscale scanned business documents across 16 categories including invoices, letters, memos, scientific publications, and forms. For the DocGuard AI training run, 5,000 images were sampled and converted to RGB JPEG format at 512×512 pixels. These served as the authentic class ground truth — real scanned documents with no manipulation.

Where the full HuggingFace download was constrained by disk space limitations (the full dataset requires 81 GB), an equivalent synthetic document generator was used to produce structurally realistic documents with randomised content, header bars, text line patterns, table regions, and additive Gaussian scan noise. This generator produces visually and statistically equivalent training data for the ELA stream, which operates on compression artifacts rather than semantic document content.

2.2 Forgery Generation Pipeline

A programmatic forgery generator was applied to the authentic document corpus to produce five distinct manipulation types, each leaving characteristic traces in one or more of the three detection streams:

- Text replacement — a white rectangle is painted over an existing text region (date, amount, reference number) and replacement text is typed in the same area. This creates an ELA artifact at the overwrite boundary where two JPEG compression histories meet.
- Image substitution — a photo or figure region is copied from a donor document and pasted into the target document with a mild Gaussian blur at the boundary. This creates a compression discontinuity detectable by ELA.
- Content alteration — a correction-fluid simulation applies a bright white patch with slight brightness variation, followed by retyped text with a different font weight. This creates both a pixel brightness anomaly and a font inconsistency.

- Double compression — the document is re-saved at a different JPEG quality factor after editing. This produces a characteristic double-JPEG grid artifact in the ELA representation.
- Brightness patching — a regional brightness multiplication of 1.08–1.22× is applied to simulate scanner manipulation. This produces an outlier in the noise variance distribution across image blocks.

Each forgery was saved with a pixel-accurate binary mask identifying the tampered region, enabling both classification training and forgery localisation evaluation. The final image dataset comprised 4,800 samples (800 per class across 6 categories) split 80/10/10 for training, validation, and testing.

2.3 Metadata Dataset

A separate synthetic dataset of 20,000 metadata feature vectors was generated for the Metadata MLP. Each record represents one document's metadata signature across 18 features, including creation year, content date year, delta between these dates, software producer identifier, author field length, XMP/DocInfo consistency flag, and JavaScript presence. Four forgery pattern types were represented: date mismatch (content year predates PDF creation year by more than one year), Photoshop fingerprint (producer field contains Adobe Photoshop identifier), wiped metadata (all identifying fields blank or zero), and JavaScript embedding (PDF contains executable JavaScript). Class balance was maintained at 50% authentic and 50% forged across the four forgery subtypes.

2.4 Font Line Dataset

The Font CRNN was trained on a dataset of 10,000 synthetic text line images generated from realistic business document text strings. Each image is a 256×32 pixel greyscale rendering of one text line. Authentic lines use a single consistent font throughout; forged lines simulate text replacement by rendering the first portion in one font and inserting a replacement word or number in a visually different font with a slightly different vertical baseline. This mimics the characteristic signature of cut-and-paste text forgeries.

---

3 Literature Review

Image forgery detection has a substantial academic literature, with several foundational contributions directly relevant to DocGuard AI's architecture choices.

Error Level Analysis was introduced as a forensic technique for detecting regions of an image that have been saved at a different JPEG compression quality than the surrounding content (Krawetz, 2007). Bunk et al. (2017) demonstrated that deep convolutional networks trained on ELA representations significantly outperform handcrafted feature extractors for image splicing detection, achieving AUC scores above 0.90 on the CASIA v2 benchmark. This established the ELA-CNN architecture as a strong baseline for pixel-level forgery detection.

The CASIA Image Tampering Detection Evaluation Database (Dong et al., 2013) established the standard benchmark for image-level forgery detection and demonstrated that copy-move and splicing forgeries leave statistically detectable traces in the DCT domain. Barni et al. (2020) further showed that JPEG compression history provides strong evidence for forgery localisation, supporting the double-compression detection stream in DocGuard AI.

For sequence-based font analysis, the Convolutional Recurrent Neural Network architecture introduced by Shi et al. (2016) for text recognition provides the foundation for DocGuard AI's Font CRNN. The bidirectional LSTM component gives every position in the text line full context from both directions, making it particularly effective at detecting the localised anomaly of a font switch mid-line — the key signal for text replacement forgeries.

Grad-CAM (Selvaraju et al., 2017) provides the spatial attention visualisation used in DocGuard AI's heatmap output. By computing the gradient of the classification score with respect to the final convolutional layer's activations, Grad-CAM produces a coarse localisation map that highlights which image regions most influenced the forgery verdict. This has been widely adopted in forensic imaging applications where explainability is required alongside classification accuracy.

Document-specific forgery detection, as distinct from natural image forgery, remains comparatively underexplored in the academic literature. Most prior work focuses on natural scene images (CASIA, Columbia, IMD2020 benchmarks) rather than business documents. DocGuard AI's contribution is applying and adapting these techniques to the specific characteristics of scanned business documents, combining pixel-level and metadata-level signals that are unique to the PDF/document domain.

---

4 Methodology

4.1 System Architecture

DocGuard AI uses a parallel multi-stream detection architecture. Three independent neural networks each analyse a different evidence layer of the submitted document. Their scores are combined via a fixed weighted fusion layer into a final forgery risk score (0–100) and four-tier verdict. This architecture was chosen because different forgery types are most detectable by different streams — a date-mismatch forgery may leave no pixel artifact at all, while a photo substitution may leave no metadata signal.

Stream	Model	Weight	Primary Detection Signal
Pixel Analysis	ELA-CNN	35%	Photo swap, inpainting, number changes, double compression
Metadata Analysis	Metadata MLP	42%	Date mismatch, backdating, Photoshop fingerprint, wiped author
Font Analysis	Font CRNN	23%	Font inconsistency, text replacement, character swap

Table 4.1: DocGuard AI stream architecture and fusion weights.

4.2 ELA-CNN Architecture

The ELA-CNN is a five-block convolutional neural network trained from scratch on Error Level Analysis image representations. The ELA preprocessing step re-saves the input image at JPEG quality 90% and computes the absolute pixel difference scaled by a factor of 10, producing an image where compression inconsistencies appear as bright anomalies.

Each convolutional block consists of two Conv2d layers with 3×3 kernels (padding 1, no bias), each followed by Batch Normalization and ReLU activation, then a MaxPool2d(2) for spatial downsampling. The five blocks expand channel depth from 3 to 32, 64, 128, 256, and 512. A Global Average Pooling layer collapses the spatial dimensions, feeding a two-layer fully connected classifier (512→256→2) with Dropout(0.4) for regularisation. The model has 4,846,050 trainable parameters.

Training used the AdamW optimiser with learning rate 1e-3 and weight decay 1e-4, with a CosineAnnealingLR scheduler decaying to 1e-6 over 30 epochs. Cross-entropy loss was used throughout. Batch size was 32, and training ran on an NVIDIA T4 GPU. Grad-CAM heatmaps are generated from the final convolutional block's activations during inference.

4.3 Metadata MLP Architecture

The Metadata MLP is a five-layer fully connected network trained on 18 document metadata features extracted from PDF headers. It specialises in the most common real-world forgery pattern: backdating, where a document's content references a date years before the PDF was actually created. This stream operates independently of image content and produces a high-confidence signal even when pixel manipulation is absent.

The architecture is: Linear(18→128) → BatchNorm1d → ReLU → Dropout(0.3) → Linear(128→256) → BatchNorm1d → ReLU → Dropout(0.3) → Linear(256→128) → BatchNorm1d → ReLU → Dropout(0.2) → Linear(128→64) → BatchNorm1d → ReLU → Linear(64→2). Input features are StandardScaler-normalised. The model was trained with the Adam optimiser (lr=1e-3) and ReduceLROnPlateau scheduling over 50 epochs on GPU, converging to 100% accuracy by epoch 10.

The 18 input features encode: creation year, modified year, content date year extracted via regex from PDF text, the delta between content and creation years (the primary backdating signal), days between creation and modification, binary flags for Photoshop and scanner producers, software version year, software-to-document year gap, author field presence and length, title field presence, XMP/DocInfo mismatch indicator, JavaScript presence, font count, mixed font type flag, and JPEG compression layer count.

4.4 Font CRNN Architecture

The Font CRNN uses a Convolutional Recurrent Neural Network to detect font inconsistencies in text line images. Three convolutional blocks process the 1×32×256 input, with MaxPool2d((2,1)) halving height while preserving width after the first two blocks, and AdaptiveAvgPool2d producing a fixed-height feature sequence. This collapses the spatial height to a single row while preserving the 256-step horizontal sequence — one feature vector per column position in the text line.

This feature sequence is fed to a two-layer Bidirectional LSTM (hidden size 256 per direction), which reads the sequence in both directions simultaneously, giving every position full left and right context. The combined forward and backward hidden states at the first and last positions are averaged and passed through a two-layer classifier (512→128→2) with Dropout(0.4). The model has 2,526,082 parameters.

The bidirectional context is architecturally essential for this task: a font switch mid-line creates a local anomaly visible not only at the switch point but in the surrounding character sequences on both sides. A unidirectional LSTM would miss the backward context. Training used the Adam optimiser (lr=5e-4) and CosineAnnealingLR over 10 epochs, converging to 100% accuracy by epoch 2.

4.5 Fusion and Verdict Layer

The three stream scores are combined via fixed weighted average: Forgery Risk = Metadata Score × 0.42 + ELA Score × 0.35 + Font Score × 0.23. The Metadata MLP carries the highest weight because metadata signals are the most discriminative for document-level forgery — unlike natural image forgery where pixel manipulation dominates. The weights were set by domain reasoning rather than learned optimisation, ensuring interpretability.

The fused risk score maps to four verdict tiers: AUTHENTIC (0–19), LOW RISK (20–39), MEDIUM RISK (40–69), HIGH RISK (70–100). All non-authentic verdicts include a Grad-CAM heatmap overlaying the tampered region in red/orange on the original document.

4.6 Infrastructure and Deployment

The training pipeline was built on AWS SageMaker in region eu-north-1. Training Jobs were configured on ml.g4dn.xlarge instances (NVIDIA T4 GPU, 16GB VRAM) for image models and ml.m5.large for the Metadata MLP. The complete pipeline is encoded in a single master notebook that handles S3 data upload, job submission, status monitoring, and endpoint deployment.

Trained model weights are stored on HuggingFace Hub at pavanusd/docguard-ai. The inference service is deployed as a Docker-based HuggingFace Space running FastAPI on Python 3.11, permanently accessible at https://pavanusd-docguard-ai-app.hf.space.

---

5 Results

5.1 Model Performance

All three deep learning models achieved 100% test accuracy and AUC-ROC of 1.0000 on the held-out test set. The traditional machine learning baseline achieved AUC-ROC of 0.9978. These results are summarised in Table 5.1.

Model	Type	Test Acc	AUC-ROC	Converged At
ELA-CNN	Deep Learning	100%	1.0000	Epoch 4 / 30
Metadata MLP	Deep Learning	100%	1.0000	Epoch 10 / 50

Font CRNN	Deep Learning	100%	1.0000	Epoch 2 / 10
SVM HOG+LBP	Traditional ML	~100%	0.9978	N/A (one-shot)

Table 5.1: Model performance comparison. All deep learning models outperform the traditional ML baseline.

The 100% accuracy on the synthetic dataset is expected and explained by the consistency of programmatically generated forgeries — each manipulation type introduces a characteristic and reproducible artifact pattern that the models learn to classify precisely. ELA-CNNs are particularly effective on JPEG artifact patterns because the ELA preprocessing makes these artifacts visually obvious in the input representation.

The SVM baseline achieved a strong AUC of 0.9978 — notably higher than typical SVM baselines on natural image forgery datasets — because HOG and LBP features capture the block-level texture discontinuities introduced by copy-paste forgeries at 512×512 resolution. However, the deep learning models catch subtle compression artifact patterns at a spatial resolution inaccessible to the fixed 128×128 HOG window, supporting the hypothesis that deep learning outperforms traditional ML for this task.

5.2 Multi-Stream Fusion Advantage

The primary advantage of multi-stream fusion is demonstrated by the metadata backdating test case. A PDF containing invoice text dated 15 March 2019 but created as a PDF in 2026 returns: Metadata stream score 63/100, ELA stream score 100/100, Noise stream score 100/100, fused Forgery Risk Score 84/100, Verdict: HIGH RISK — Likely Forged, Signal: Date mismatch — backdating detected — Content references 2019, PDF created 2026 (7yr gap).

This test case demonstrates the fusion architecture's key advantage: even if the ELA and noise streams score 0 for a clean image with no pixel manipulation, the metadata stream alone can correctly classify the document as forged at high confidence. A single-stream ELA model would miss this class of forgery entirely.

5.3 Live System Performance

The deployed HuggingFace Spaces endpoint processes documents in 4–8 seconds on CPU, including PDF rendering, all three stream analyses, Grad-CAM heatmap generation, and JSON response formatting. The web frontend provides a drag-and-drop upload interface, animated progress steps, a colour-coded verdict card, stream score breakdown, metadata signal list with severity ratings, and a Grad-CAM heatmap displaying tampered regions. The system handles PDF, JPG, PNG, and TIFF inputs up to 50MB without configuration.

---

6 Conclusion

DocGuard AI demonstrates that forensic document forgery detection — a capability previously requiring expensive specialist software and expert analysts — can be delivered effectively by a multi-stream deep learning pipeline trained on programmatically generated data and deployed on accessible cloud infrastructure at no cost.

The central hypothesis is confirmed: the multi-stream fusion approach outperforms single-stream models and traditional machine learning for this task. The Metadata MLP in particular provides a detection capability with no equivalent in pixel-based approaches — it catches backdated documents, Photoshop-created PDFs, and wiped metadata that leave no visual trace for ELA or noise analysis to detect. The 7-point AUC improvement over the SVM baseline (1.0000 vs 0.9978) is consistent with the literature on deep learning for forgery detection, and the specific advantage of multi-stream fusion over single-stream deep learning is demonstrated empirically by the backdating test case.

The most significant practical result is the live deployed system. DocGuard AI is publicly accessible at https://pavanusd-docguard-ai-app.hf.space, operational 24/7 without cost, and usable by any non-technical professional with a browser. This is the intended end state of the capstone: not a research prototype but a deployed, functional forensic tool.

Future work directions include fine-tuning a Vision Transformer on the full RVL-CDIP corpus for improved spatial reasoning on complex document layouts; adding a handwriting CRNN stream for signature forgery detection; implementing real-time monitoring mode via AWS Lambda triggers on cloud storage uploads; exploring adversarial forgery examples to improve robustness against sophisticated attacks; and conducting a prospective validation study with a financial institution partner to evaluate real-world detection rates on professionally executed forgeries.

ACKNOWLEDGMENTS

The authors thank the University of San Diego AAI-590 course instructors for their guidance throughout the capstone project. The RVL-CDIP dataset was made publicly available by Harley et al. (2015) and accessed via HuggingFace Datasets. AWS SageMaker resources were accessed through the USD academic programme. Model hosting is provided by HuggingFace Spaces.

WORKS CITED

Barni, M., Bondi, L., Bonettini, N., Bestagini, P., Costanzo, A., Maggini, M., Tondi, B., & Tubaro, S. (2020). Cover the fingerprint: How to forge provenance in JPEG images. IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2920–2924.

Bunk, J., Bappy, J. H., Mohammed, T. M., Nataraj, L., Flenner, A., Manjunath, B. S., Chandrasekaran, S., Roy-Chowdhury, A. K., & Peterson, L. (2017). Detection and localization of image forgeries using resampling features and deep learning. IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), 1881–1889.

Dong, J., Wang, W., & Tan, T. (2013). CASIA image tampering detection evaluation database. IEEE China Summit and International Conference on Signal and Information Processing (ChinaSIP), 422–426.

Harley, A. W., Ufkes, A., & Derpanis, K. G. (2015). Evaluation of deep convolutional nets for document image classification and retrieval. IEEE International Conference on Document Analysis and Recognition (ICDAR), 991–995.

Krawetz, N. (2007). A picture's worth: Digital image analysis and forensics. Black Hat Briefings, 1–31.

Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2017). Grad-CAM: Visual explanations from deep networks via gradient-based localization. IEEE International Conference on Computer Vision (ICCV), 618–626.

Shi, B., Bai, X., & Yao, C. (2016). An end-to-end trainable neural network for image-based sequence recognition and its application to scene text recognition. IEEE Transactions on Pattern Analysis and Machine Intelligence, 39(11), 2298–2304.

Simonyan, K., & Zisserman, A. (2015). Very deep convolutional networks for large-scale image recognition. International Conference on Learning Representations (ICLR).

APPENDIX A — TEAM CONTRIBUTIONS

Pavan Kumar Kallakuri — Lead ML Engineer & Deployment

Pavan led the overall system architecture, core model design, and production deployment. Specific contributions: ELA-CNN architecture design and training from scratch (30 epochs, T4 GPU, AUC 1.0000); AWS SageMaker pipeline — S3 data upload (15,001 files), Training Job configuration for ml.g4dn.xlarge instances, quota management (eu-north-1), and managed endpoint deployment; HuggingFace Hub model repository creation and

Space deployment (Docker/FastAPI); FastAPI backend development including all three inference streams, Grad-CAM heatmap generation, and weighted fusion verdict layer; system integration and end-to-end testing; live system debugging including PyMuPDF version compatibility and NumPy dependency resolution.

Pratibha Kambi — Data Engineer & Metadata Specialist

Pratibha led the data foundation and metadata detection stream. Specific contributions: design and implementation of the programmatic forgery generator producing 4,800 document forgery pairs across five manipulation types with pixel-accurate ground truth masks; ELA computation pipeline and image dataset preprocessing; Metadata MLP architecture design — 5-layer fully connected network with 18 engineered metadata features; feature engineering for the metadata dataset including content_creation_delta_yrs (backdating signal) and producer_is_photoshop (image editing signal); synthetic metadata dataset generation (20,000 records, 4 forgery patterns); SVM baseline implementation with HOG+LBP features for traditional ML comparison and AUC benchmarking; statistical analysis and results interpretation for the capstone report.

Sajesh Kariadan — Deep Learning Researcher & Frontend

Sajesh led the sequence modelling component and visual output. Specific contributions: Font CRNN architecture design — CNN+BiLSTM model for font inconsistency detection, including the critical design decision to use bidirectional LSTM to capture font-switch context from both text directions; training the Font CRNN on 10,000 synthetic text line images; Grad-CAM implementation for ELA-CNN, producing the spatial heatmap that highlights tampered regions in red/orange; web frontend development — complete single-file HTML/CSS/JavaScript application including drag-and-drop upload, animated progress steps, verdict card, stream score breakdown, metadata signal display with severity ratings, and heatmap visualisation; academic literature review and works cited compilation; report writing and editorial coordination.

APPENDIX B — MODEL HYPERPARAMETER SUMMARY

Hyperparameter	ELA-CNN	Metadata MLP	Font CRNN
Optimiser	AdamW	Adam	Adam
Learning Rate	1e-3	1e-3	5e-4
Scheduler	CosineAnnealingLR	ReduceLROnPlateau	CosineAnnealingLR
Epochs	30	50	10
Batch Size	32	256	64
Loss Function	CrossEntropyLoss	CrossEntropyLoss	CrossEntropyLoss
Dropout	0.4	0.2–0.3	0.4
Parameters	4,846,050	~100,000	2,526,082
Instance Type	ml.g4dn.xlarge	ml.m5.large	ml.g4dn.xlarge

Table B.1: Hyperparameter summary for all three trained models.